

SQL + DATA MODELING

Artist names are not identifiers.

How often does one exact name refer to
several MusicBrainz artist entities?

2.9M

**artist entities
queried in PostgreSQL**

MusicBrainz full export, 2026-07-15

Two counts, two different meanings.

126,415

duplicate exact names

Each row of the grouped result represents one name.


419,770

artist entities affected

This is the sum of all entities behind those duplicate names.

COUNT(*) answers the first question. SUM(entity_count) answers the second.

One name can hide many different entities.



Indigo: 249 artist entities across 52 non-null areas.

The SQL is simple. The grain is the real lesson.

```
WITH name_counts AS (  
    SELECT name, COUNT(*) AS entity_count  
    FROM artist  
    GROUP BY name  
    HAVING COUNT(*) > 1  
)  
SELECT COUNT(*), SUM(entity_count)  
FROM name_counts;
```

A name is an attribute.
A MusicBrainz ID is an entity key.

Joining on names can silently connect unrelated entities.

What this analysis does not prove.

- Exact primary names only
- Aliases are excluded
- [unknown] and [no artist] excluded from the ranking
- Catalog coverage is not popularity

NEXT EPISODE

What happens when artist names are used as join keys?

MusicBrainz SQL Stories #2