

4,320

TRACK ROWS FOR ONE RECORDING

**I THOUGHT THIS
RECORDING
APPEARED
EVERYWHERE.**

The query looked enormous.

4,320

TRACK ROWS

BUT ONLY

1

RELEASE

My SQL was counting correctly.
My first interpretation was not.

RECORDING

I Feel Bad For Your Hard Drive

Release ID 5055218

ONE RELEASE CONTAINS THE WHOLE RESULT.



**180 MEDIA × 24 MATCHING
TRACK ROWS = 4,320**

release -> medium -> track

Each yellow marker is drawn from the checked 180 × 24 structure.

FIRST, CHANGE THE RESULT GRAIN.

```
WITH recording_usage AS (  
    SELECT track.recording AS recording_id,  
           COUNT(track.id) AS track_appearances  
    FROM track  
    GROUP BY track.recording  
)
```

ONE RESULT ROW = ONE RECORDING

COUNT(track.id) answers:
How many track rows?

It does not answer:
Across how many releases?

```
track -> medium -> release  
COUNT(DISTINCT release.id)
```

The second count checks the interpretation, not the SQL syntax.

A CORRECT COUNT CAN STILL SUPPORT A WRONG INTERPRETATION.

Before interpreting an aggregate,
state what one row represents
and what the count actually measures.

MusicBrainz full export / 2026-07-15

Catalog metadata, not plays, sales,
audience, or popularity.

github.com/leodid68/musicbrainz-sql-stories